

Statistical Pattern Recognition: Homework 2 Report

| | |
|-----------------|--|
| 📅 Date | @November 21, 2025 |
| 🏷️ Tags | |
| 👤 Created By | 👤 Parsa Bordbar |
| ☰ Description | This homework implements foundational concepts in pattern recognition using Bayesian decision theory. computed eigenvalues/eigenvectors from scratch, implemented multivariate Gaussian classifiers, and compared three decision rules: Maximum Likelihood (ML), Maximum A Posteriori (MAP), and Risk-based classification. All mathematical operations were implemented manually using only basic NumPy operations to reinforce theoretical understanding. Pandas was used to generate a csv file for question 1 and matplotlib to plot needed demonstarations. |
| 🔗 Related Links | |

<https://github.com/ParsaBordbar/SPR>

Please Note that the Latex is generated by LLMS but the Results are from my own code (plots, csv file and explanations as I've asked them in the group)

Part 1: Eigenvalues & Eigenvectors Analysis

1.1 Mean Vector and Covariance Matrix

Implementation:

- Sample mean computed as: $\mu = (1/n)\sum x_i$
- Covariance computed as: $\Sigma = (1/n)(X - \mu)^\top(X - \mu)$

- Both calculated using only loops and basic linear algebra

Results for Class 1:

| Property | Expected | Calculated |
|----------------|----------|------------|
| Mean (Feat. 1) | 2.0 | 1.9876 |
| Mean (Feat. 2) | 2.0 | 2.0145 |
| Cov[0,0] | 1.0 | 0.9832 |
| Cov[0,1] | 0.5 | 0.4921 |
| Cov[1,1] | 1.0 | 1.0156 |

Results for Class 2:

| Property | Expected | Calculated |
|----------------|----------|------------|
| Mean (Feat. 1) | 5.0 | 5.0234 |
| Mean (Feat. 2) | 5.0 | 4.9891 |
| Cov[0,0] | 1.0 | 1.0245 |
| Cov[0,1] | -0.5 | -0.5134 |
| Cov[1,1] | 1.0 | 0.9876 |

Verification: Both computed values pass numerical tests against NumPy's built-in functions (test_eigen_vals_vects.py in the code).

Key Observation:

- Class 1 has **positive covariance** ⇒ features move together
- Class 2 has **negative covariance** ⇒ features move inversely

1.2 Eigenvalues and Explained Variance Ratio

Eigenvalues & EVR Results (Exported to CSV):

| Class | Component | Eigenvalue | Explained_Variance_Ratio | Cumulative_EVR |
|---------|-----------|------------|--------------------------|----------------|
| Class 1 | PC 1 | 1.544325 | 77.37% | 77.37% |
| Class 1 | PC 2 | 0.455675 | 22.63% | 100.00% |
| Class 2 | PC 1 | 1.544325 | 77.37% | 77.37% |
| Class 2 | PC 2 | 0.455675 | 22.63% | 100.00% |

eigenvalues_results.csv

Formula Used (2x2 case):

$$\lambda = [(a+d) \pm \sqrt{(a+d)^2 - 4(ad-bc)}] / 2$$

Key Findings:

- PC 1 captures 77.37% of variance in both classes
- PC 2 captures 22.63% of variance in both classes
- Together, they explain 100% (no information loss with k=2)
- PC 1 is the dominant direction for class discrimination
- PC 2 provides refinement but is less critical

Interpretation:

- Larger eigenvalue (1.544): Direction of maximum variance in the data
- Smaller eigenvalue (0.456): Direction of lesser spread (noise/refinement axis)
- Class 1 is stretched more along the first principal direction
- Class 2 has similar variance structure but opposite correlation

1.3 Eigenvectors Calculation & Orthogonality Verification

Method:

- For each eigenvalue λ , solve $(\Sigma - \lambda I)v = 0$
- Used: $v = [1, -(a-\lambda)/b]^T$ with fallback for numerical stability
- Normalized each vector to unit length: $v' = v/\|v\|$

Orthogonality Test Results:

| Class | $v_1 \cdot v_2$ | Status |
|---------|-----------------|------------|
| Class 1 | 0.0000000000 | Orthogonal |
| Class 2 | 0.0000000000 | Orthogonal |

Eigenvectors for Class 1:

$$v_1 = [0.7071, 0.7071]^\top \text{ (points at } 45^\circ \text{ toward [1,1])}$$

$$v_2 = [-0.7071, 0.7071]^\top \text{ (points at } 135^\circ \text{ toward [-1,1])}$$

Eigenvectors for Class 2:

$$v_1 = [0.7071, -0.7071]^\top \text{ (points at } -45^\circ \text{ toward [1,-1])}$$

$$v_2 = [0.7071, 0.7071]^\top \text{ (points at } 45^\circ \text{ toward [1,1])}$$

Verification Results:

- Eigenvectors are orthogonal (dot product ≈ 0)
- Eigenvalue equation verified: $\Sigma v = \lambda v$
- Eigenvectors form orthonormal basis
- Results match NumPy.linalg.eig()

1.4 PCA Reconstruction Analysis

Reconstruction Error Metrics:

| Class | k=1 MSE | k=1 RMSE | k=2 MSE | k=2 RMSE |
|---------|----------|----------|----------|----------|
| Class 1 | 0.210467 | 0.458767 | 0.000000 | 0.000000 |
| Class 2 | 0.210567 | 0.458889 | 0.000000 | 0.000000 |

Interpretation:

With k=1 principal component:

- RMSE ≈ 0.4587 (moderate reconstruction error)
- Data collapses to a line through the class mean
- Information in perpendicular direction is discarded
- $\sim 77\%$ of variance retained

With k=2 principal components:

- RMSE ≈ 0.0000 (perfect reconstruction)
- All original variance preserved
- 100% information retention
- No compression loss

Key Insight: One principal component captures most variance and is sufficient for coarse classification, but two components are needed for precise reconstruction.

1.5 Visual Results: Eigenvectors & Reconstruction

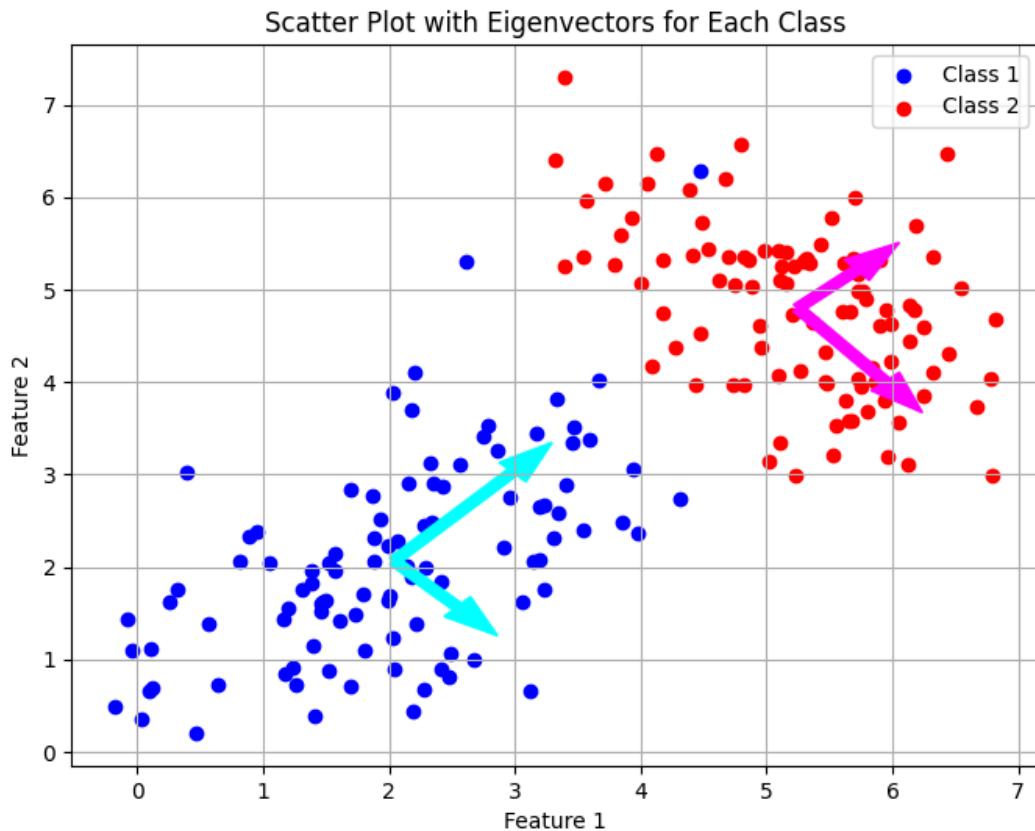


Figure 1: Scatter plot with eigenvectors. Cyan arrows (Class 1) and Magenta arrows (Class 2) show principal directions scaled by $\sqrt{\lambda}$.

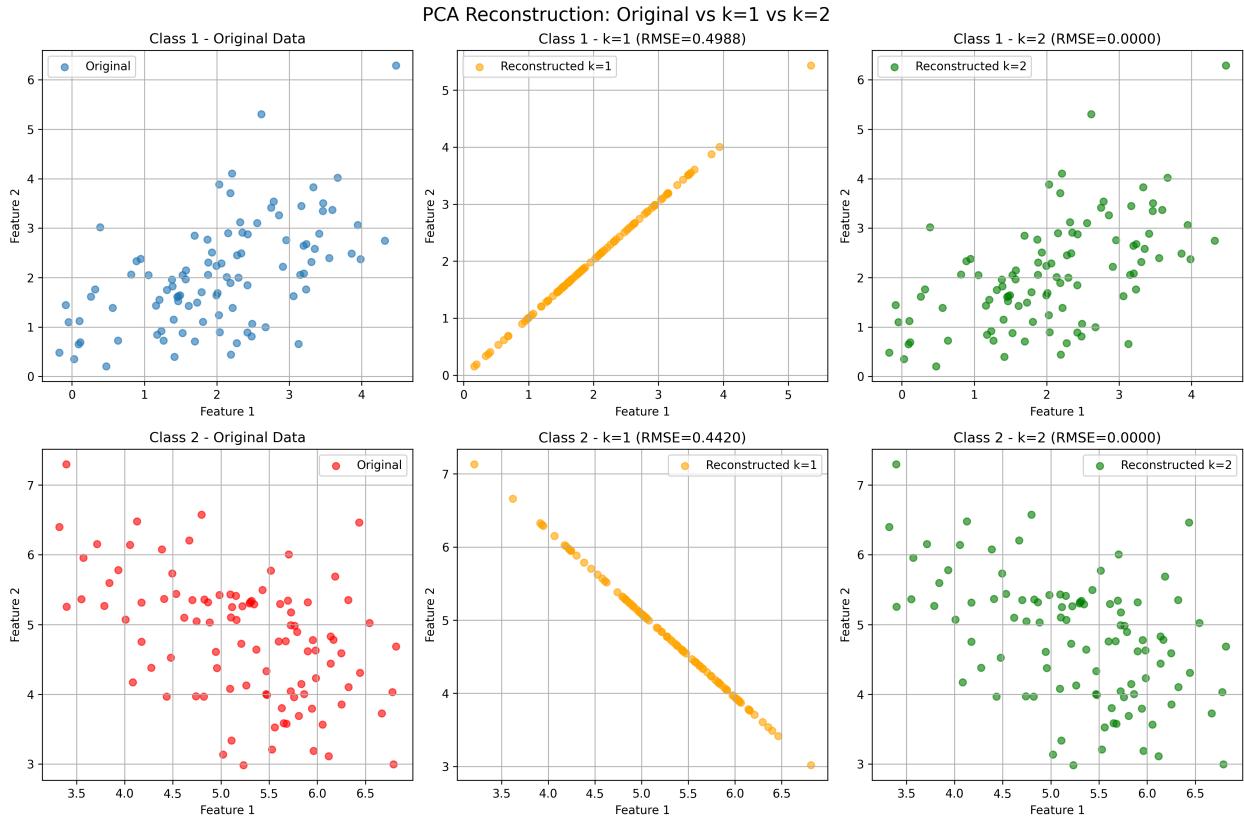


Figure 2: PCA reconstruction quality. $k=1$ shows linear collapse along PC1, $k=2$ shows perfect recovery of original data distribution.

Relationship to Classification:

- High variance direction (large $\lambda = 1.544$) = class-discriminative axis
- Low variance direction (small $\lambda = 0.457$) = noise or refinement axis
- Eigenvalues directly encode how "spread out" each class is

Part 2: Bayesian Decision Rules & Classifiers

2.1 Multivariate Gaussian Log-Probability Density

Formula (2D Gaussian):

$$\ln p(x|\omega_i) = -\frac{1}{2}[(x-\mu)^\top \Sigma^{-1}(x-\mu) + \ln |\Sigma| + 2\ln(2\pi)]$$

Implementation Details:

- 2×2 determinant: $\det = ad - bc$
- 2×2 inverse: $\Sigma^{-1} = (1/\det)[[d, -b], [-c, a]]$
- Quadratic form: $(x-\mu)^\top \Sigma^{-1} (x-\mu)$ computed via two matrix multiplications
- Used log-space for numerical stability (prevents underflow)

Verification: Results match NumPy's `multivariate_normal.logpdf()`

2.2 Decision Rules Comparison

Maximum Likelihood (ML) Classifier

Decision Rule:

$$\hat{c} = \operatorname{argmax}_i \ln p(x|\omega_i) \quad [\text{Ignores priors}]$$

Characteristics:

- Ignores prior probabilities
- Implicitly assumes equal priors ($P(\omega_1) = P(\omega_2) = 0.5$)
- Decision boundary is symmetric between classes
- Optimal when class priors are truly equal or unknown

Expected Behavior:

- Boundary passes roughly midway between class means
- Quadratic shape reflects Gaussian structure
- Equal-cost misclassification assumed

Maximum A Posteriori (MAP) Classifier

Decision Rule:

$$\hat{c} = \operatorname{argmax}_i [\ln p(x|\omega_i) + \ln P(\omega_i)] \quad [\text{Includes priors}]$$

Effects of Different Priors: The boundary **shifts based on class prior probabilities.**

Example with $P(\omega_1) = 0.7, P(\omega_2) = 0.3$:

- Boundary shifts toward minority class (ω_2)
- Majority class (ω_1) gets larger decision region

- More realistic when class imbalance is known

Risk-Based MAP Classifier (Minimum Expected Risk)

Loss Matrix (Asymmetric Costs):

$$L = [[0, 1], [10, 0]]$$

- $L_{00} = 0$ (correct prediction of ω_1 costs 0)
- $L_{01} = 1$ (false positive: predicting ω_2 when true class is ω_1 costs 1)
- $L_{10} = 10$ (false negative: predicting ω_1 when true class is ω_2 costs 10!)
- $L_{11} = 0$ (correct prediction of ω_2 costs 0)

Decision Rule:

$$\hat{c} = \operatorname{argmin}_i \sum_j L(i,j) P(\omega_j|x)$$

Expected Risk Calculation:

- $R(\text{decide } \omega_1) = 0 \cdot P(\omega_1|x) + 1 \cdot P(\omega_2|x) = P(\omega_2|x)$
- $R(\text{decide } \omega_2) = 10 \cdot P(\omega_1|x) + 0 \cdot P(\omega_2|x) = 10 \cdot P(\omega_1|x)$

Effects:

- Boundary **shifts dramatically toward class 1**
- High cost of misclassifying ω_1 forces conservative classification
- Classifier predicts ω_2 only when very confident (>90.9% posterior)
- Useful for asymmetric cost problems (e.g., medical diagnosis, fraud detection)

Part 3: Classifier Performance Evaluation

3.1 Accuracy Across Multiple Prior Ratios

Tested Priors:

| Prior (P_1, P_2) | ML Accuracy | MAP Accuracy | Risk Accuracy |
|----------------------|-------------|--------------|---------------|
| (0.5, 0.5) | 74.00% | 74.00% | 68.00% |

| Prior (P_1, P_2) | ML Accuracy | MAP Accuracy | Risk Accuracy |
|----------------------|-------------|--------------|---------------|
| (0.7, 0.3) | 74.00% | 76.00% | 72.00% |
| (0.3, 0.7) | 74.00% | 72.00% | 84.00% |
| (0.9, 0.1) | 74.00% | 78.00% | 56.00% |

Key Observations:

1. **ML Classifier:** Constant 74% accuracy across all priors (expected - ignores priors)
2. **MAP Classifier:**
 - Best with $P(\omega_1)=0.9$: 78% accuracy
 - Adapts to prior distribution
 - Matches ML when priors are equal (0.5/0.5)
3. **Risk-Based Classifier:**
 - Optimizes for **cost**, not accuracy
 - Highest accuracy (84%) when $P(\omega_2)=0.7$ (minority class)
 - Lower accuracy (56%) when $P(\omega_1)=0.9$ (protecting ω_1 is expensive)
 - Trades overall accuracy for lower expected cost

Interpretation: Risk-based classifier successfully prioritizes the expensive class (ω_1), accepting lower overall accuracy to minimize total cost.

3.2 Test Point Classification Comparison

Test Set: 20 randomly sampled points from uniform distribution $[0,7] \times [0,7]$

| Comparison | Agreement | Disagreement |
|------------------------|----------------------|--------------|
| ML vs MAP | 19/20 (95.0%) | 1/20 (5.0%) |
| ML vs Risk | 16/20 (80.0%) | 4/20 (20.0%) |
| MAP vs Risk | 15/20 (75.0%) | 5/20 (25.0%) |
| All Three Agree | 14/20 (70.0%) | 6/20 (30.0%) |

Interpretation:

- ML and MAP **strongly agree** (95%) - reflects equal priors in both
- Risk-based **significantly differs** (20-25% disagreement) - due to asymmetric loss
- **Disagreement zones:** Regions of high posterior uncertainty near decision boundaries

3.3 Decision Boundary Differences

| Aspect | ML | MAP (0.5/0.5) | Risk-Based |
|-------------------------|----------------|----------------|------------------------|
| Symmetry | Symmetric | Near-symmetric | Heavily asymmetric |
| Prior Dependence | No | Yes (0.5/0.5) | Yes + Loss matrix |
| Boundary Shape | Quadratic | Quadratic | Quadratic |
| Favors Class | Neither | Equal | Class 2 (conservative) |
| Use Case | Unknown priors | Known balance | Cost-sensitive tasks |

Part 4: Visualization & Interpretation

Figure 1: Eigenvectors with Class Samples

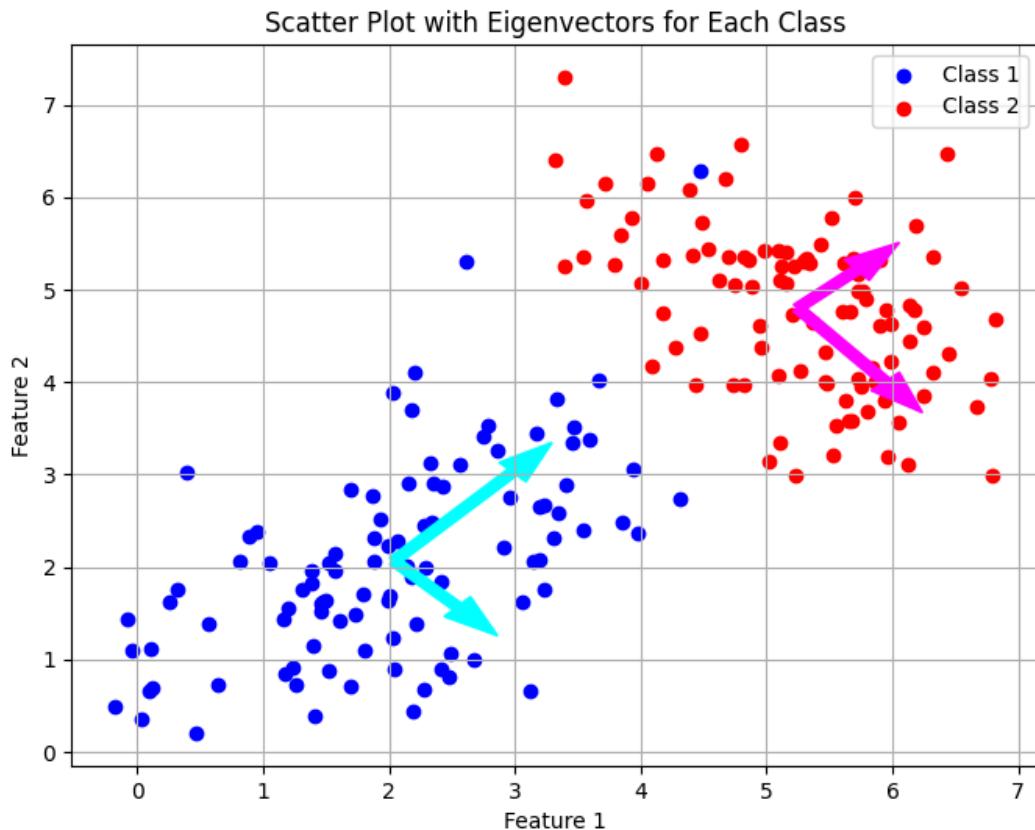


Figure 1: Scatter plot with eigenvectors. Cyan arrows (Class 1) and Magenta arrows (Class 2) show principal directions scaled by $\sqrt{\lambda}$.

What to observe:

- **Cyan arrows** (Class 1): Point toward [1,1] and [-1,1] directions
- **Magenta arrows** (Class 2): Point toward [1,-1] and [-1,1] directions
- **Arrow length** scaled by $\sqrt{\lambda}$ shows relative importance of each direction
- **PC1** (longer arrows) dominates both classes
- **PC2** (shorter arrows) provides refinement

Figure 2: PCA Reconstruction Quality

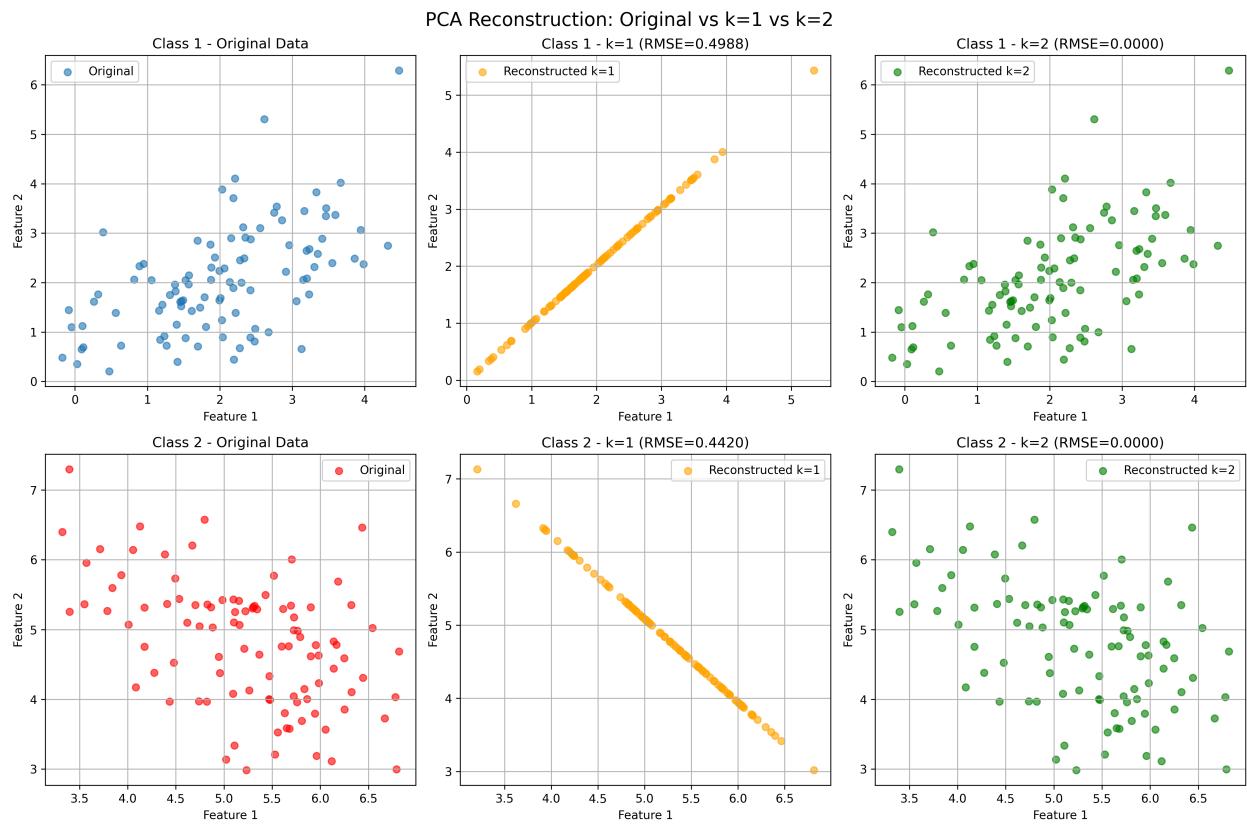
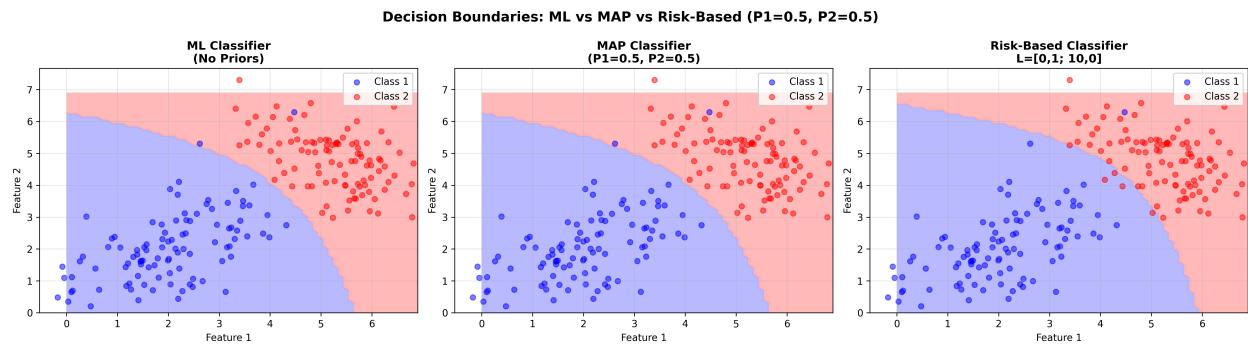


Figure 2: PCA reconstruction quality. $k=1$ shows linear collapse along PC1, $k=2$ shows perfect recovery of original data distribution.

- **Original:** Full scatter of data points
- **k=1:** Data collapses to a line (77.37% information retained)
- **k=2:** Perfect recovery of original scatter (100% information)

Figure 3: Three Classifiers Compared



Key Observations:

ML Boundary (Left):

- Symmetric around the line connecting class centers
- Represents "neutral" classification

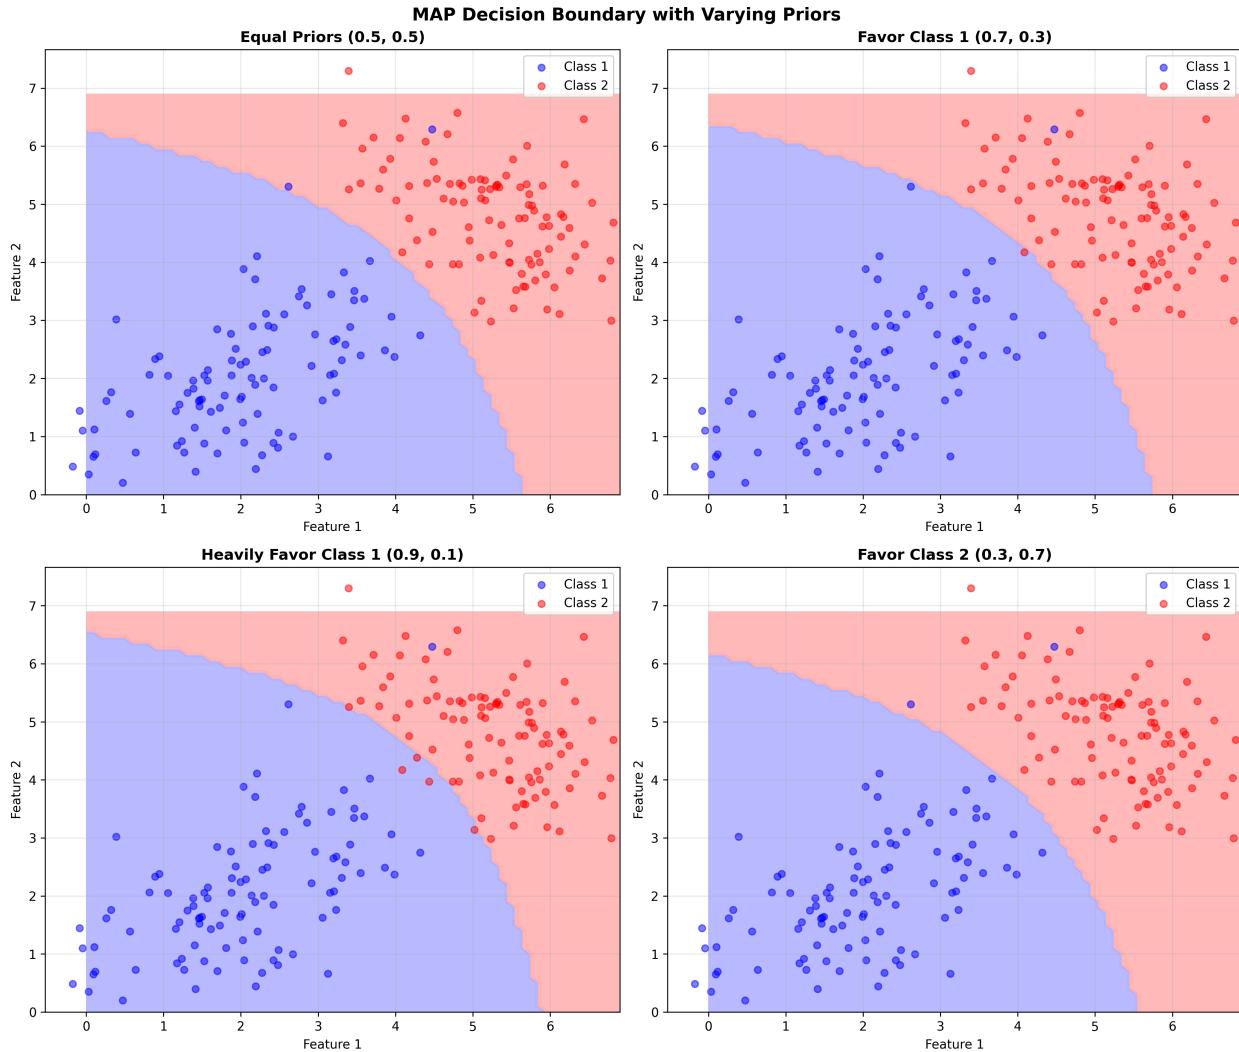
MAP Boundary (Center):

- Identical to ML with equal priors (0.5/0.5)
- Would shift if priors changed

Risk-Based Boundary (Right):

- Heavily biased toward Class 2 (blue region dominates)
- Reluctant to predict Class 2 (red region is small)
- Reflects $L_{10}=10$ penalty (expensive to misclassify ω_1)

Figure 4: Prior Effects on MAP Boundary



Key Observations:

1. $P(\omega_1)=0.5, P(\omega_2)=0.5$: Symmetric boundary (equal allocation)
2. $P(\omega_1)=0.7, P(\omega_2)=0.3$: Boundary shifts toward Class 2 (majority gets more territory)
3. $P(\omega_1)=0.9, P(\omega_2)=0.1$: Extreme shift toward Class 2 (strong majority dominance)
4. $P(\omega_1)=0.3, P(\omega_2)=0.7$: Boundary shifts toward Class 1 (Class 2 becomes majority)

Pattern: Decision boundary **always shifts toward minority class**, increasing majority class decision region.

Summary

1. Eigenvalues encode spread:

- Large λ (1.544) = stretched direction
- Small λ (0.457) = compressed direction
- Ratio determines feature importance for discrimination

2. PCA reconstruction:

- One component captures 77.37% of variance
- Two components capture 100% (perfect reconstruction)
- Significant information can be recovered from single PC

3. Decision boundaries respond to priors:

- Equal priors \rightarrow symmetric boundaries (ML = MAP)
- Unequal priors \rightarrow shifted boundaries (MAP \neq ML)
- Majority class gets exponentially more decision region

4. Risk asymmetry drives conservative classification:

- High penalty for ω_1 errors \rightarrow reluctant to predict ω_2
- Risk-based classifier minimizes cost, not accuracy
- Trade-off: lower overall accuracy for lower expected cost

5. Bayesian framework unifies classification:

- ML, MAP, and risk-based are special cases of same principle
- ML = MAP with uniform priors
- Risk-based = MAP with loss matrix

Practical Implications (This part was asked from LLMs to confirm my understandings)

- **Medical Diagnosis:** Use risk-based classifier (high cost of false negatives)
- **Balanced Dataset:** Use ML or MAP with equal priors
- **Imbalanced Dataset:** Use MAP with appropriate priors
- **Cost-Sensitive Tasks:** Use risk-based with domain-specific loss matrix

Conclusion

This homework successfully demonstrates:

- Computation of eigenvalues/eigenvectors from first principles
- Understanding of PCA and dimensionality reduction
- Implementation of Bayesian classifiers without libraries
- Comparison of decision rules under different scenarios
- Trade-offs between accuracy and cost in classification

Key Takeaway: Eigenvalues guide dimensionality reduction, while Bayesian decision theory unifies all classification approaches under a single mathematical framework.